

Quant II

Lab 2

Yinxuan Wang

2026-02-12

Homework 1

- Homework 1 due February 16 before class
- Email to **both** Cyrus & me
- Send one PDF, including math, code, and output
- Math part can be scanned and attached (or from tablet)
LaTeX recommended

Today's Agenda

- Robust standard errors
- Cluster robust standard errors
- Bootstrap

Robust Variance Estimators

- From lecture, asymptotic variance of OLS under unit-level sampling

$$V = \mathbb{E}[X_i X_i']^{-1} \mathbb{E}[X_i X_i' \varepsilon_i^2] \mathbb{E}[X_i X_i']^{-1}$$

Robust Variance Estimators

- From lecture, asymptotic variance of OLS under unit-level sampling

$$V = \mathbb{E}[X_i X_i']^{-1} \mathbb{E}[X_i X_i' \varepsilon_i^2] \mathbb{E}[X_i X_i']^{-1}$$

- A general form of estimators for the asymptotic variance:

$$\hat{V}(\cdot) = (\mathbf{X}' \mathbf{X})^{-1} \mathbf{X}' \hat{\Psi} \mathbf{X} (\mathbf{X}' \mathbf{X})^{-1}$$

- Where $\hat{\Psi}$ is an estimator of $\text{plim}[\varepsilon \varepsilon']$

Robust Variance Estimators

- From lecture, asymptotic variance of OLS under unit-level sampling

$$V = \mathbb{E}[X_i X_i']^{-1} \mathbb{E}[X_i X_i' \varepsilon_i^2] \mathbb{E}[X_i X_i']^{-1}$$

- A general form of estimators for the asymptotic variance:

$$\hat{V}(\cdot) = (\mathbf{X}' \mathbf{X})^{-1} \mathbf{X}' \hat{\Psi} \mathbf{X} (\mathbf{X}' \mathbf{X})^{-1}$$

- Where $\hat{\Psi}$ is an estimator of $\text{plim}[\varepsilon \varepsilon']$
- These estimators are “robust” because they don’t require specific distributional assumptions

Robust Variance Estimators

From *Mostly Harmless Econometrics*, Ch.8.

- HC0: $\hat{\psi}_i = \hat{e}_i^2$

Robust Variance Estimators

From *Mostly Harmless Econometrics*, Ch.8.

- HC0: $\hat{\psi}_i = \hat{e}_i^2$
- HC1: $\hat{\psi}_i = \frac{n}{n-k} \hat{e}_i^2$

Robust Variance Estimators

From *Mostly Harmless Econometrics*, Ch.8.

- HC0: $\hat{\psi}_i = \hat{e}_i^2$
- HC1: $\hat{\psi}_i = \frac{n}{n-k} \hat{e}_i^2$
- HC2: $\hat{\psi} = \frac{1}{1-h_{ii}} \hat{e}_i^2$, where h_{ii} is the leverage of unit i (the influence of i on the regression line)

Robust Variance Estimators

From *Mostly Harmless Econometrics*, Ch.8.

- HC0: $\hat{\psi}_i = \hat{e}_i^2$
- HC1: $\hat{\psi}_i = \frac{n}{n-k} \hat{e}_i^2$
- HC2: $\hat{\psi} = \frac{1}{1-h_{ii}} \hat{e}_i^2$, where h_{ii} is the leverage of unit i (the influence of i on the regression line)
- HC3: $\hat{\psi} = \frac{1}{(1-h_{ii})^2} \hat{e}_i^2$

Linear Regression in R

- In base R, `lm()` implements the OLS estimation
 - Fit the model
 - Then use functions in the package `sandwich` for robust SEs

Linear Regression in R

- In base R, `lm()` implements the OLS estimation
 - Fit the model
 - Then use functions in the package `sandwich` for robust SEs
- With packages `estimatr`, `lfe`, or `fixest`
 - Specify the SE estimator from inside the function

Linear Regression in R

- In base R, `lm()` implements the OLS estimation
 - Fit the model
 - Then use functions in the package `sandwich` for robust SEs
- With packages `estimatr`, `lfe`, or `fixest`
 - Specify the SE estimator from inside the function
- Check the function documentation to see what SEs are computed

Robust Standard Errors in R, i

```
set.seed(123)
pacman::p_load(tidyverse, estimatr, sandwich, lmtest)

# simulate population
pop <- data.frame(Y1 = rnorm(1000, 4, 2), Y0 = rnorm(1000, 0.5, 3))

# random sample
sample <- pop[sample(nrow(pop), 100),]
sample$D <- 0
sample$D[sample(100, 30)] <- 1

# observed potential outcomes
sample <- sample %>% mutate(Y = D*Y1 + (1-D)*Y0)
```

Robust Standard Errors in R, ii

```
### Example 1: lm_robust  
# Default: HC2 estimator  
lm_robust(Y ~ D, data = sample)
```

	Estimate	Std. Error	t value	Pr(> t)	CI Lower	CI Upper
(Intercept)	0.3137636	0.3844703	0.8160932	4.164261e-01	-0.4492052	1.076732
D	4.1641379	0.5135960	8.1078083	1.493054e-12	3.1449234	5.183352

```
# HC1
```

```
lm_robust(Y ~ D, data = sample, se_type = "HC1")
```

	Estimate	Std. Error	t value	Pr(> t)	CI Lower	CI Upper
(Intercept)	0.3137636	0.3855896	0.8137243	4.177757e-01	-0.4514264	1.078954
D	4.1641379	0.5128987	8.1188318	1.414176e-12	3.1463072	5.181969

```
# or
```

```
lm_robust(Y ~ D, data = sample, se_type = "stata")
```

	Estimate	Std. Error	t value	Pr(> t)	CI Lower	CI Upper
(Intercept)	0.3137636	0.3855896	0.8137243	4.177757e-01	-0.4514264	1.078954
D	4.1641379	0.5128987	8.1188318	1.414176e-12	3.1463072	5.181969

Robust Standard Errors in R, iii

```
### Example 2: lm + lmtest + sandwich
fit <- lm(Y ~ D, data = sample)
summary(fit) # Default is homoskedastic
```

Call:

```
lm(formula = Y ~ D, data = sample)
```

Residuals:

Min	1Q	Median	3Q	Max
-5.7960	-2.1130	-0.0758	2.4649	6.0472

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.3138	0.3446	0.910	0.365
D	4.1641	0.6292	6.618	1.95e-09 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.884 on 98 degrees of freedom

Multiple R-squared: 0.3089, Adjusted R-squared: 0.3018

F-statistic: 43.79 on 1 and 98 DF, p-value: 1.949e-09

```
N <- nrow(sample)
D <- cbind(rep(1, N), sample$D)
K <- dim(D)[2]
vcov <- solve((t(D) %*% D)) %*% t(D) %*% diag((sum(residuals(fit)^2)/(N-K)), N, N) %*% D %*% solve((t(D) %*% D)
cbind(coef(fit), sqrt(diag(vcov)))
```

	[,1]	[,2]
(Intercept)	0.3137636	0.3446480
D	4.1641379	0.6292383

Robust Standard Errors in R, iv

```
coeftest(fit, vcov = vcovHC(fit, type = "HC1"))
```

t test of coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.31376	0.38559	0.8137	0.4178
D	4.16414	0.51290	8.1188	1.414e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
coeftest(fit, vcov = vcovHC(fit, type = "HC2"))
```

t test of coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.31376	0.38447	0.8161	0.4164
D	4.16414	0.51360	8.1078	1.493e-12 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Export Regression Output in Tables - stargazer

```
pacman::p_load(stargazer)
stargazer(fit, type = "text")
```

Dependent variable:

Y

D	4.164***
	(0.629)
Constant	0.314
	(0.345)
Observations	100
R2	0.309
Adjusted R2	0.302
Residual Std. Error	2.884 (df = 98)

Export Regression Output in Tables - modelsummary

```
pacman::p_load(modelsummary)
modelsummary(list(fit, fit),
            vcov=list("HC1", "HC2"),
            output = "markdown", gof_omit = "^(!R2|Num)")
```

	(1)	(2)
(Intercept)	0.314 (0.386)	0.314 (0.384)
D	4.164 (0.513)	4.164 (0.514)
Num.Obs.	100	100
R2	0.309	0.309
R2 Adj.	0.302	0.302

Export Regression Output in Tables - fixest::etable

```
pacman::p_load(fixest)
fix <- list(feols(Y ~ D, sample, vcov="iid"),
            feols(Y ~ D, sample, vcov="HC1"))
fix %>% etable()
```

	model 1	model 2
Dependent Var.:	Y	Y
Constant	0.3138 (0.3446)	0.3138 (0.3856)
D	4.164*** (0.6292)	4.164*** (0.5129)

S.E. type	IID	Heteroskeda.-rob.
Observations	100	100
R2	0.30886	0.30886
Adj. R2	0.30181	0.30181

Signif. codes:	0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1	

Clustering

- Clustered **sampling**
 - For instance, we sample villages or electoral precincts: individual behavior (e.g. voting) is correlated within cluster

Clustering

- Clustered **sampling**
 - For instance, we sample villages or electoral precincts: individual behavior (e.g. voting) is correlated within cluster
- Clustered **treatment assignment**
 - For instance, we randomize a treatment at the housing block level, assignment is correlated within cluster

Clustering

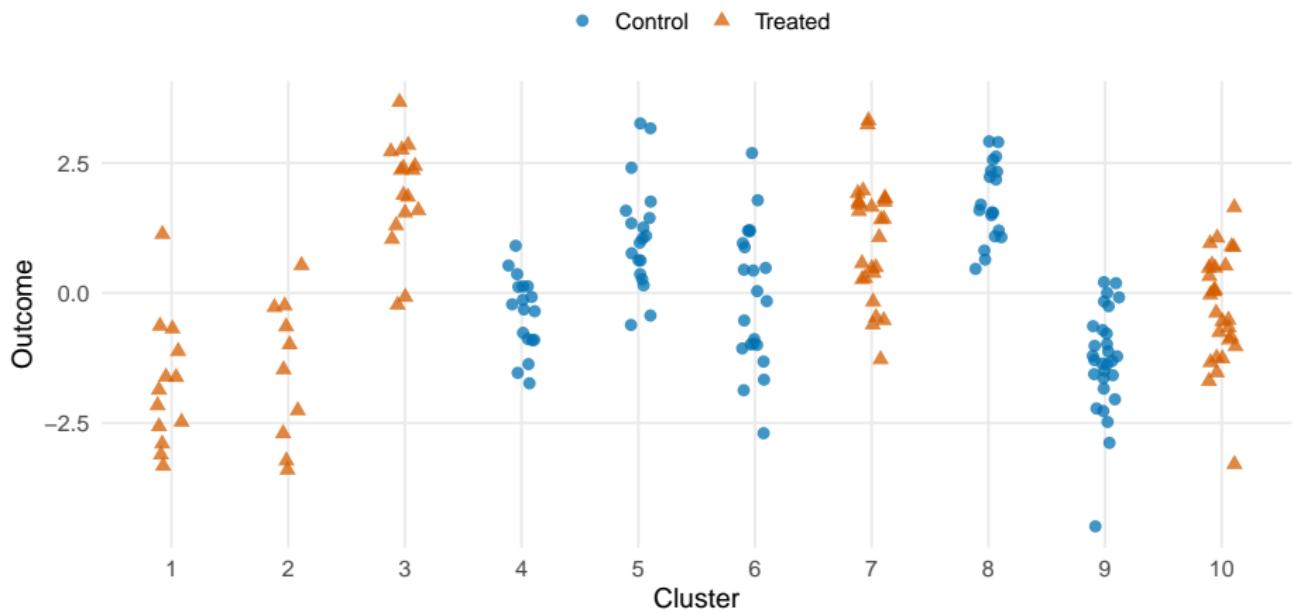
- Clustered **sampling**
 - For instance, we sample villages or electoral precincts: individual behavior (e.g. voting) is correlated within cluster
- Clustered **treatment assignment**
 - For instance, we randomize a treatment at the housing block level, assignment is correlated within cluster
- In both cases, relative to simple random process, units are contributing less information
- If we treat units as independent
 - we overstate the information in the data
 - and underestimate the variance of the estimator

Level of Clustering

A guide to empirical practice - MacKinnon, Nielsen, and Webb (*JE*, 2023)

- In observational work we do not perfectly observe the cluster structure
- “How do you cluster SEs?”
- Traditional rules of thumb:
 - When in doubt, cluster at the highest possible level
 - Choose the level which gives the highest SEs
- DiD setting: cluster at the unit level
 - Can also cluster in two dimensions if suspects of correlation along other dimensions
- If treatment is assigned at a higher level: cluster at that level
 - E.g. individual-level data, geographic distances at place level

Cluster Simulation

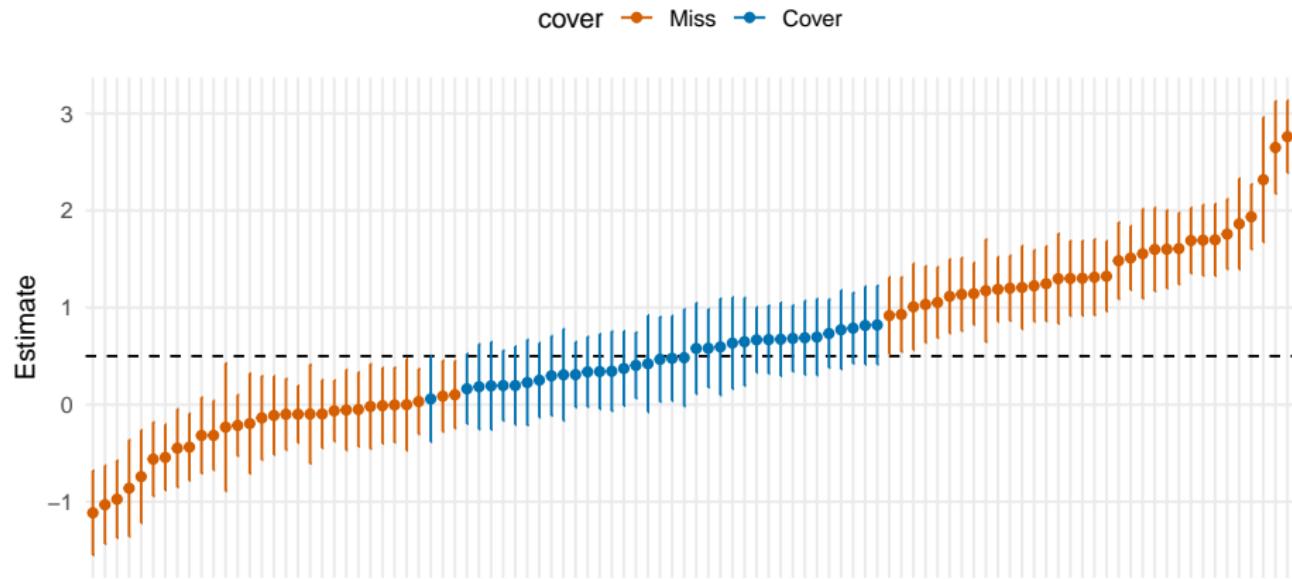


Cluster Simulation, ii

```
set.seed(123)
make_sim <- function(x) {
  df <- make_clustered_data(G = 10, N = 200)
  model <- summary(lm(y ~ treat, data = df))$coefficients["treat", ]
  model
}
res <- map(1:100, make_sim) %>%
  bind_rows() %>%
  mutate(
    lo = Estimate - 1.96 * `Std. Error`,
    hi = Estimate + 1.96 * `Std. Error`,
    cover = lo < 0.5 & hi > 0.5,
    i = 1:100
  )
```

Cluster Simulation, iii

- Only 36% cover the true value



Cluster Simulation - Clustered SEs

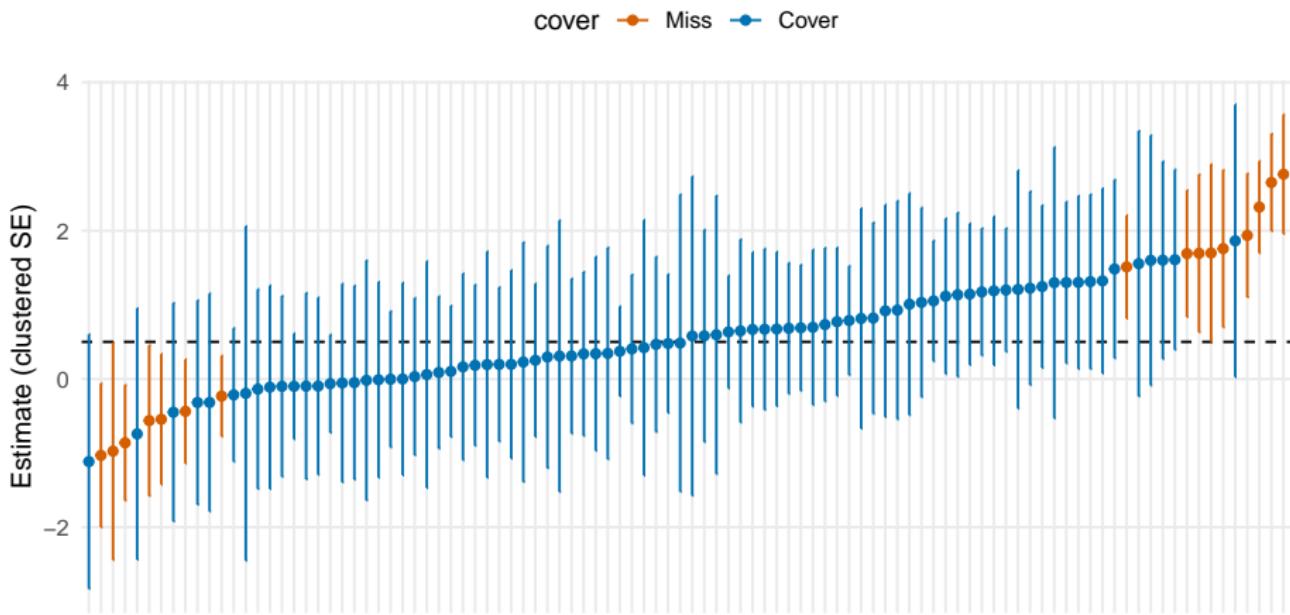
- lm + sandwich
- lfe::felm
- fixest::feols

```
set.seed(123)
make_sim_clustered <- function(x) {
  df <- make_clustered_data(G = 10, N = 200)
  m <- feols(y ~ treat, data = df, cluster = ~ g)

  ct <- coeftable(m)
  out <- ct["treat", c("Estimate", "Std. Error")]
  tibble(Estimate = out[1], `Std. Error` = out[2])
}
```

Cluster Simulation - Clustered SEs, ii

- 84% cover the true value



Bootstrapping

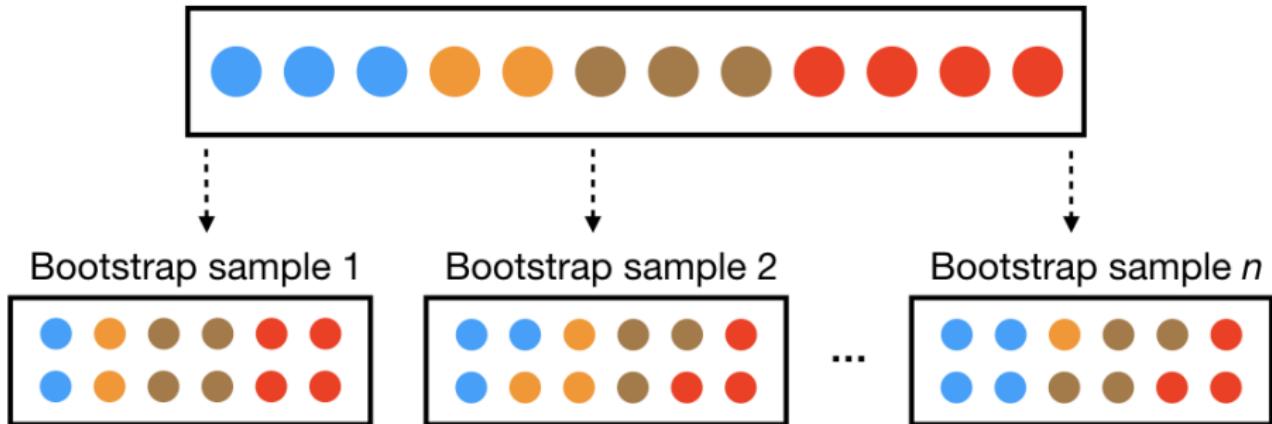
- Simulate sampling distributions (from repeated random sampling) using the available sample
- Parametric, Semi-parametric, Non-parametric
- **Non-parametric:** based on re-sampling
- Random re-sampling from the sample approximates random sampling from the population
- The sampling distribution of statistics from re-samples approximates the true sampling distribution

Vanilla Non-parametric Bootstrapping

- ① From our sample of size N , draw a random sample of size N with replacement
- ② On this sample, compute the estimate
- ③ Repeat many times (e.g. 1000)
- ④ Obtain a distribution of estimates from the resamples (a bootstrapped sampling distribution)
- ⑤ For inference, compute moments of the bootstrapped distribution

Vanilla Non-parametric Bootstrapping

Original Data



Cluster Simulation, Bootstrapping

```
set.seed(123)
df <- make_clustered_data(G = 30, N = 500)
n_clusters <- length(unique(df$g))
```

Cluster Simulation, Vanilla Unit-level Bootstrap

Let's try to bootstrap SEs for β_{treat}

Cluster Simulation, Vanilla Unit-level Bootstrap

Let's try to bootstrap SEs for β_{treat}

```
vanilla_bootstrap <- function(a) {  
  resampled.data <- df %>% sample_frac(1, replace = T)  
  broom::tidy(lm(y ~ treat, data = resampled.data)) %>%  
    dplyr::filter(term == "treat")  
}  
vb.results <- map(1:1000, vanilla_bootstrap) %>%  
  bind_rows()  
sd(vb.results$estimate)
```

```
[1] 0.1325565
```

Cluster Simulation, Vanilla Unit-level Bootstrap

Let's try to bootstrap SEs for β_{treat}

```
vanilla_bootstrap <- function(a) {
  resampled.data <- df %>% sample_frac(1, replace = T)
  broom::tidy(lm(y ~ treat, data = resampled.data)) %>%
    dplyr::filter(term == "treat")
}
vb.results <- map(1:1000, vanilla_bootstrap) %>%
  bind_rows()
sd(vb.results$estimate)
```

```
[1] 0.1325565
```

```
tidy(summary(lm(y ~ treat, data = df)))
```

```
# A tibble: 2 x 5
  term      estimate std.error statistic   p.value
  <chr>        <dbl>     <dbl>     <dbl>      <dbl>
1 (Intercept) -0.0102     0.114    -0.0900  0.928
2 treat         0.560      0.141     3.98    0.0000799
```

Cluster Simulation, Vanilla Cluster Bootstrap

- Bootstrap procedure with clustered data needs to reflect the clustered structure
- Let's resample at the **cluster** level

Cluster Simulation, Vanilla Cluster Bootstrap

- Bootstrap procedure with clustered data needs to reflect the clustered structure
- Let's resample at the **cluster** level

```
get_cluster_bs_sample <- function(df) {  
  clusters <- unique(df$g)  
  sampled_clusters <- sample(clusters, size = n_clusters, replace = TRUE)  
  dplyr::bind_rows(lapply(sampled_clusters, function(cl) df[df$g == cl, ]))  
}  
  
estimates <- purrr::map_dbl(1:1000, function(b) {  
  boot_df <- get_cluster_bs_sample(df)  
  coef(lm(y ~ treat, data = boot_df))["treat"]  
})  
  
sd(estimates)  
  
[1] 0.4648703  
tidy(summary(feols(y ~ treat, data = df, cluster =~ g)))  
  
# A tibble: 2 x 5  
  term      estimate std.error statistic p.value  
  <chr>     <dbl>     <dbl>     <dbl>    <dbl>  
1 (Intercept) -0.0102     0.353    -0.0290   0.977  
2 treat        0.560      0.459     1.22     0.232
```

Wild Cluster Bootstrap- t

Let's not draw conventional bootstrap samples

Instead,

- Fit the unrestricted model, estimate $\hat{\beta}$ and t

Wild Cluster Bootstrap- t

Let's not draw conventional bootstrap samples

Instead,

- Fit the unrestricted model, estimate $\hat{\beta}$ and t
- Fit the restricted model, obtain residuals

Wild Cluster Bootstrap- t

Let's not draw conventional bootstrap samples

Instead,

- Fit the unrestricted model, estimate $\hat{\beta}$ and t
- Fit the restricted model, obtain residuals
- **Randomize the sign of the residual at the cluster level**
(Rademacher weights)

Wild Cluster Bootstrap- t

Let's not draw conventional bootstrap samples

Instead,

- Fit the unrestricted model, estimate $\hat{\beta}$ and t
- Fit the restricted model, obtain residuals
- **Randomize the sign of the residual at the cluster level**
(Rademacher weights)
- Construct Y_b^* as restricted fitted values + randomized residuals

Wild Cluster Bootstrap- t

Let's not draw conventional bootstrap samples

Instead,

- Fit the unrestricted model, estimate $\hat{\beta}$ and t
- Fit the restricted model, obtain residuals
- **Randomize the sign of the residual at the cluster level**
(Rademacher weights)
- Construct Y_b^* as restricted fitted values + randomized residuals
- Estimate $\hat{\beta}_b$

Wild Cluster Bootstrap- t

Let's not draw conventional bootstrap samples

Instead,

- Fit the unrestricted model, estimate $\hat{\beta}$ and t
- Fit the restricted model, obtain residuals
- **Randomize the sign of the residual at the cluster level**
(Rademacher weights)
- Construct Y_b^* as restricted fitted values + randomized residuals
- Estimate $\hat{\beta}_b$
- Compute t -statistic, t_b , for the new estimate

Wild Cluster Bootstrap- t

Let's not draw conventional bootstrap samples

Instead,

- Fit the unrestricted model, estimate $\hat{\beta}$ and t
- Fit the restricted model, obtain residuals
- **Randomize the sign of the residual at the cluster level**
(Rademacher weights)
- Construct Y_b^* as restricted fitted values + randomized residuals
- Estimate $\hat{\beta}_b$
- Compute t -statistic, t_b , for the new estimate
- Do this N_{boot} times

Wild Cluster Bootstrap- t

Let's not draw conventional bootstrap samples

Instead,

- Fit the unrestricted model, estimate $\hat{\beta}$ and t
- Fit the restricted model, obtain residuals
- **Randomize the sign of the residual at the cluster level**
(Rademacher weights)
- Construct Y_b^* as restricted fitted values + randomized residuals
- Estimate $\hat{\beta}_b$
- Compute t -statistic, t_b , for the new estimate
- Do this N_{boot} times
- Compute bootstrap p-values by counting the share of simulated t_b to the left/right of the observed t

Cluster Simulation, Wild Cluster Bootstrap-t

```
pacman::p_load(fwildclusterboot)
# install.packages(c("dqrng", "collapse", "JuliaConnectoR", "gtools"))

feols_fit <- feols(y ~ treat, data = df)
boot <- fwildclusterboot::boottest(
  feols_fit,
  clustid = "g",
  param   = "treat",
  B       = 9999,
  type    = "rademacher",
  conf_int = TRUE,
  sign_level = 0.05
)
summary(boot)

boottest.fixest(object = feols_fit, param = "treat", B = 9999,
  clustid = "g", sign_level = 0.05, conf_int = TRUE, type = "rademacher")

Hypothesis: 1*treat = 0
Observations: 500
Bootstr. Type: rademacher
Clustering: 1-way
Confidence Sets: 95%
Number of Clusters: 30

      term estimate statistic p.value conf.low conf.high
1 1*treat = 0     0.56      1.22   0.253   -0.416     1.588
```

R Packages for Bootstrapping

- `boot`
- `ClusterBootstrap` - bootstrapping on hierarchically structured data
- `fwildclusterboot` - wild cluster bootstrap
- `multiwaycov::cluster.boot` for post-estimation SE calculation